


Cryptography

One-Time Pad: Two-Pad Attack

- Reuse of same key stream leaks relation between plaintexts: $C_1 = P_1 \oplus K, C_2 = P_2 \oplus K \Rightarrow C_1 \oplus C_2 = P_1 \oplus P_2$
- If one plaintext is known or guessable, the other can be recovered.
- Equal plaintexts produce equal ciphertexts (pattern leak).
- Rule: Never reuse key stream (also applies to stream ciphers).


Stream vs Block Ciphers

- Generate pseudo-random key stream $S = \text{StreamCipher}(IV, K)$.
- Encryption: $C = P \oplus S$.
- Fast, bit/byte oriented, low latency.
- Fatal if nonce reused → same two-pad attack.



Block Ciphers

- Operate on fixed block size (AES: 128-bit).
- Modes of operation semantics: ECB – deterministic, pattern leaks, avoid.
- CBC – uses random IV, needs MAC for integrity. CTR – turns block cipher into stream, needs unique nonce. GCM – AEAD mode (confidentiality + integrity).
- Pros: hardware support, flexible AEAD modes.
- Cons: padding issues (CBC), IV misuse.



Theoretical Properties

Symmetric Encryption

- Provides confidentiality.
- Security defined by indistinguishability (IND-CPA / IND-CCA).
- Key space 2^k , brute force cost $\approx 2^{(k-1)}$.

Asymmetric (Public-Key)

- Uses key pair (pk, sk).
- Enables key exchange, encryption without pre-shared secret, and signatures.
- Slower; usually used only for key establishment.
- Hash Functions**
- Map variable input to fixed n-bit output.
- Properties: Preimage resistance – hard to find input for given hash. Second-preimage resistance – hard to find different input with same hash. Collision resistance – hard to find any two inputs with same hash ($\approx 2^{(n/2)}$ effort).
- Used for: integrity, digital signatures, password hashing, HMAC, commitments.

Problem Solved by Asymmetric Cryptography

- Addresses the key distribution problem: allows secure communication without pre-shared secret.
- Anyone can encrypt using public key; only private key holder can decrypt.
- Enables digital signatures, hybrid encryption (Use PK share symmetric key).

Confidentiality, Integrity, Authentication

- Confidentiality
- Only authorized parties can read the message.
- Achieved by encryption (e.g., AES, ChaCha20).

Integrity

- Message cannot be modified undetected.
- Achieved by MAC, HMAC, AEAD, or digital signatures.

Authentication

- Entity authentication: verify who is communicating (e.g., challenge-response).
- Message authentication: verify origin and integrity of message (MAC or signature).


Relationships

- Encryption alone → confidentiality only.
- Signature or MAC → integrity + authentication.
- AEAD (AES-GCM, ChaCha20-Poly1305) → confidentiality + integrity together.
- Signatures do not hide message content.

Digital Signatures

- Security Properties
- Authenticity – verifies identity of signer (via certified public key).
- Integrity – any change in message breaks verification.
- Non-repudiation – signer cannot deny valid signature.

Structure



- Alice gets identified by the CA
- Alice's public key K_A^{pub} gets signed by the (trustworthy) CA using the CA's private key $K_{\text{CA}}^{\text{pri}}$
- The signed public key K_A^{pub} is linked to Alice and the CA provides a certificate
- Anyone that wants to confirm the certificate can check the CA's signature by encrypting the certificate using the CA's public key $K_{\text{CA}}^{\text{pub}}$

Cryptographic Implementations:


Performance parameters for hardware design

- Area (silicon used), Throughput (data per time), Power (peak), Energy (per task), Time-to-Market (design time), Latency (time per item)

Important parameters when implementing crypto in HW/SW

- From HW perspective: fast speed (sometimes), low latency, low area, low power, low energy.
- From SW perspective / benchmarking: know the platform (word size, available instructions), measure clock cycles, memory occupation, power, energy; aim for fair, reproducible results.

Public-Key Encryption / Decryption



Software implementation: LUTs vs compute-on-the-fly

- LUT (tables):** implement S-boxes and even whole round functions as look-ups – fewer instructions, higher speed; costs more memory. Best when code footprint and cache allow.

- Compute on the fly:** cheaper in memory, portable across platforms; may increase cycles (consider shift/rotate costs on target CPU). Balance using platform knowledge (word size, instruction costs).

- Practical approach:** profile on the target, report cycles and memory; pick the mix (partial tables + computed steps) that minimizes cycles under your memory budget.

Crypto algorithm selection nowadays (what guides the choice)

- Start from specifications: algorithm standard and system requirements drive the design flow.
- Then: Use standardized, publicly scrutinized algorithms (e.g., AES, RSA/ECC), their match to constraints (area, power, latency, throughput).
- Trend facts from slides: RSA (large keys) is slower; many systems moved to ECC; post-quantum migration is underway (lattice-based). Choose based on required security level and platform performance.

Cryptanalysis

What is a fault attack / DFA (high level)

- Attacker injects faults during computation (active attacker), obtains a fault-free ciphertext C and a faulted ciphertext C' for the same plaintext

- P, and uses $C \oplus C'$ (and how the fault propagates) to recover key material (divide-and-conquer on S-boxes / round parts).

Countermeasures (fault attacks)

- Redundancy / repetition: run operation twice and compare outputs (check mismatch → alarm).
- Error detection / correction for internal states.
- Sensors for voltage/clock/laser anomalies.
- Note:** redundancy or comparisons themselves can be faulted (so use hardened checks / multi-layer defenses).

Timing attacks — idea & countermeasures


- Idea (simple):** If execution time depends on secret data (branches, memory/cache behavior), attacker measures time (locally or remotely via RTT) and infers secret bits (e.g., PIN checking with early exit reveals prefix bits).

Practical countermeasures (timing)

- Constant-time implementations: remove data-dependent branches and memory access patterns.
- Avoid LUTs or ensure accesses do not depend on secret (compute S-box algebraically instead of table lookup; or use platform features).
- Hardware fixes:** special hardware that enforces constant-time memory access (slower).
- Random delays:** can increase noise but do not guarantee security (CLT reduces effectiveness with many samples).


Side channel attacks:

- AES rounds**



10 repeating patterns thus it's AES-128, last round smaller due to the lack of MixColumns

RSA square-and-multiply leak



Modular exponentiation does square every loop and multiply only when exponent bit = 1. Presence/absence of the multiply shows up as longer execution and different power patterns (Timing+Power pattern attack).


Correlation Power Analysis (CPA) Attack

Step 0: To perform CPA we must know the cipher's structure

- All values (p, k, x, y) have size of 8 bits

$$\triangleright p = k \oplus k$$

$$\triangleright y = S(x)$$



Step 1: Choose an intermediate value v of the cipher to attack

v must be a function of the plaintext and the key i.e. $v = f(p, k)$

A common choice for v is the Sbox output i.e. $v = y = S(p \oplus k)$


Note that recovering y is equivalent to recovering k since: $x = S^{-1}(y)$, the Sbox S is one-to-one and p is known, thus $k = x \oplus p$

Step 2: Measure the power consumption of the device

- In the measurement setup we generate randomly n 8-bit plaintexts p_i
- Typically the number of traces n is large i.e. thousands to millions
- Store the plaintexts in the input vector $p = [p_1, p_2, p_3, \dots, p_n]$
- While encrypting every p_i , digitized signal over time i.e. the trace t_i
- Every trace t_i contains m time samples i.e.


$$t_i = [t_i[1], t_i[2], t_i[3], \dots, t_i[m]]$$

- Capturing n power traces with m time samples per trace results in the measurement matrix M
- At some point in time the device should process the intermediate $v = y = S(p \oplus k_{\text{true}})$ that we target




Step 3: Compute the hypothetical intermediate values

- In our device $v = y = S(p \oplus k_{\text{true}})$, but k_{true} is unknown
- For any plaintext p_i we can compute the value y_i for all possible key candidates $k \in \{0, 1, \dots, 255\}$, only 1 key (1 col) correct



Step 4: Apply the leakage model

- We map the hypothetical intermediate values to hypothetical power consumption values, producing the power-prediction matrix
- A common choice for leakage model is **Hamming weight (HW)** of a value



Step 5: Compare power-prediction matrix PP with measurement matrix M

- Use Pearson's correlation coefficient: uncorrelated columns $\rho_{a,b} \approx 0$, correlated columns $\rho_{a,b} \approx \pm 1$, CPA use absolute correlation $|\rho_{a,b}|$
- Compare columns of power-prediction PP with measurement M

$$\rho_{a,b} = \frac{\sum_{i=1}^n \left(a_i - \bar{a} \right) \left(b_i - \bar{b} \right)}{\sqrt{\sum_{i=1}^n \left(a_i - \bar{a} \right)^2} \sqrt{\sum_{i=1}^n \left(b_i - \bar{b} \right)^2}}$$

Hamming weight (HW) and Hamming distance (HD) models

- Hamming weight (HW)** of value v: number of '1' bits in v. Example: $HW(10011)=3$.
- HW often models bus precharge or software register loads (microcontrollers) where power correlates with number of ones.

- **Hamming distance (HD)** between two values a,b: number of bit flips to go from a to b ($HD(a,b) = HW(a \oplus b)$). Example: $HD(10011,01101) = 4$.
- HD models dynamic switching (register transitions) and is common for hardware (FPGA/ASIC).
- **Use in attacks:** Map hypothesized intermediate values to HW/HD predictions and correlate with measured traces (see CPA next).

Intrusion Detection Systems:

Confusion Matrix and Basic Metrics

- True Positive $TP = P(\text{alert}'|\text{intrusion}') = P(A|I)$
- True Negative $TN = P(\text{'no alert}'|\text{'no intrusion}') = P(\neg A|\neg I)$
- False Positive $FP = P(\text{'alert}'|\text{'no intrusion}') = P(A|\neg I)$
- False Negative $FN = P(\text{'no alert}'|\text{'intrusion}') = P(\neg A|I)$

Reality		Intrusion	No Intrusion
IDS Decision	Intrusion	True Positive	False Positive
	No Intrusion	False Negative	True Negative

- Base rate parameter $B = P(I)$ i.e. the probability of an intrusion
- Bayesian detection rate (positive predictive value) PPV
 - PPV close to 1: after we see an alert, since an intrusion is likely
 - PPV close to 0: can ignore the alert, since an intrusion is unlikely

$$PPV = P(I|A) = \frac{P(A,I)}{P(A)} = \frac{P(A|I)P(I)}{P(A)} =$$

$$\frac{P(A|I)P(I)}{\sum_{X \in \{I, \neg I\}} P(A,X)} = \frac{P(A|I)P(I)}{\sum_{X \in \{I, \neg I\}} P(A|X)P(X)} = \frac{P(A|I)P(I)}{P(A|I)P(I) + P(A|\neg I)P(\neg I)} \iff \\ PPV = \frac{TP * B}{TP * B + FP * (1 - B)}$$

- Similarly, negative predictive value $NPV = P(\neg I|\neg A)$


Base Rate Fallacy (PPV can be low even with good TP)

- With rare intrusions (small B), even small FP dominates alerts.
- Example: $TP=0.99$, $FP=0.01$, $B=0.001 \rightarrow PPV \approx 0.99-0.001/(0.99-0.001+0.01-0.999) \approx 0.09$ (~9% of alerts are real).

- Takeaway: Good IDS must good at suppressing false alerts, low FP


Decision Trees: Roll-Back (Expected-Cost) Analysis

- Leaves to Parent node, best operating point for IDS minimizes cost



Receiver Operator Characteristic (ROC) Curve

- 2d graph showing tradeoff between **TP** (y-axis) and **FP** (x-axis) in [0, 1]



- **low TP / low FP** 'cautious': will raise an alert if strong evidence

- **high TP / high FP** 'relaxed': will raise an alert if weak evidence

- Goal is high TP low FP, above the line better than random guess

Distance Based Outliers (DBO)

- Definition: A datapoint o in D is a DB(p, d) outlier if at least fraction p of the datapoints in D lie in distance greater than d from o

Construction

Step 1: Construct the following set S:

$S = \{x \in D \text{ s.t. } dist(o, x) \leq d\}$ $dist(o, x) = \text{Euclidean distance between } o, x$

Step 2: Compute the set cardinality $|S|$

If $|S| \leq (1-p)|D|$ then the datapoint o is an outlier

If $|S| > (1-p)|D|$ then the datapoint o is normal behavior

Simplify the anomaly detection problem and use distance metric

Isolation Tree (Isolation Forest)

Construction

Assume that we have the following training dataset with:

► 6 datapoints stored in vectors $\{x_0, x_1, x_2, x_3, x_4, x_5\}$


► 5 features $\{x_0, x_1, x_2, x_3, x_4\}$ for each datapoint xi

$$D = \begin{bmatrix} x_0^T \\ x_1^T \\ x_2^T \\ x_3^T \\ x_4^T \\ x_5^T \end{bmatrix} = \begin{bmatrix} 4.74 & 0.71 & -0.26 & 2.79 & 5.11 \\ 4.71 & 0.75 & -0.30 & 2.50 & 5.58 \\ 4.72 & 0.82 & -0.28 & 2.61 & 5.26 \\ 4.66 & 0.72 & -0.31 & 2.62 & 5.30 \\ 4.80 & 0.71 & -0.27 & 2.71 & 5.33 \\ 4.81 & 0.77 & -0.29 & 2.72 & 5.40 \end{bmatrix}$$

Step 1: Randomly select a feature q from $\{x_0, x_1, x_2, x_3, x_4\}$ e.g. $q = x_2$

Step 2: Randomly select a split point p between max and min of q e.g. $\max(x_2) = -0.26$ and $\min(x_2) = -0.31$ $p = [-0.31, -0.26]$, pick $p = -0.28$

Step 3: Partition dataset D to the left and right sets D_L and D_R such that: $D_L = \{x \in D \text{ s.t. } q < p\}$ and $D_R = \{x \in D \text{ s.t. } q > p\}$ refer as **isolation tree** (iTTree)



Step 4: The construction of the iTTree continues recursively

▪ **iForest idea:** On average, the path length of anomalies is shorter compared to the path length of normal datapoints

RNG, IoT & OS security

Secure OS: Remove all the unnecessary software; Remove all the unnecessary services; Change default accounts ; Use the least privilege principle; Regularly update; Activate logging

Valid for IoT devices?

▪ Partly valid, but IoT is different: No antivirus; Limited support lifetime; Poor user interface; The programmers are not battle-hardened; The computer is connected to a physical device.

▪ Emphasize signed firmware, secure boot, automatic updates/rollback, strong device-hub/vendor auth, network isolation (separate VLAN), default-credential removal, and "hub as firewall/IDS" patterns.

Why IoT protection different from PCs/servers:

- Limited interfaces (no keyboard/screen) → weak configuration & visibility
- No/limited antivirus model and subscription economics
- Constrained resources → constant-time crypto and lightweight stacks; avoid heavy agents.
- Attached to physical processes → safety requires sanity checks (hardware interlocks).
- Cloud/vendor reliance → vendor compromise/abandonment risks; signed firmware a must.
- Usability economics → users rarely update; defaults persist; recovery/reset needs physical controls.

Pseudo Random Number Generator (PRNG)

- **Definition:** state $\leftarrow F(\text{state}, \text{seed})$ Generated from a seed; Sequences that have the statistical property of a random sequence
- **Useful in security?** A normal PRNG is deterministic and predictable once its seed or internal state is known → not safe for keys or nonces.

True random number generator

[Entropy Source] → [Digitizer / Sampler] → [Post-Processing / Extractor] → [Ring Oscillators] → [Quantum]

User authentication

Password entropy H(X)

$$H(X) = - \sum_{x \in X} p(x) \log_2 p(x)$$

Entropy is expressed in bits

A higher number of bits implies more uncertainty

When computing the entropy $H(X)$ discard zero probability

It holds that $H(X) \geq 0$ always

Dictionary attacks and Salt values

▪ **Verification:** store (user, h) , where $h = \text{hash}(\text{password})$. At login, compute $h' = \text{hash}(\text{input})$ and compare $h' == h$.

▪ **Method:** Let dictionary D that contains b most common passwords $D = \{x_1, x_2, \dots, x_b\}$ Attacker can hash them in advance, i.e. create a lookup table

$$H = \{\text{hash}(x_1), \text{hash}(x_2), \dots, \text{hash}(x_b)\}$$

Attack amounts to a single table lookup in H thus it is much faster than computing hashes on-the-fly

▪ **Salt value:** a non-secret value for each password

Compute: $h = \text{hash}(x||\text{salt})$, Store in database: $(\text{userid}, h, \text{salt})$

To verify a user password x' Compute: $h' = \text{hash}(x'||\text{salt})$, Compare: $h == h'$

▪ **Two-factor authentication (2FA):** To login now I need something that I know + something that I have.

- e.g. a password together with a passcode generator / a password together with a hardware security token / a password together with mobile authenticator app / a PIN together with a smartcard

Privacy & Anonymity

k-anonymity

Let a database D with non-sensitive attributes $\{A_1, A_2, \dots, A_n\}$

Let a quasi-identifier $Q = \{A_i, \dots, A_j\} \subset \{A_1, A_2, \dots, A_n\}$

Let the database projection $D[Q]$

i.e. discard all attributes in D that are not in the Q set

k-Anonymity. We say that the database D satisfies k-anonymity with respect to quasi-identifier Q iff every sequence of values in the projection $D[Q]$ appears with at least k occurrences

Benefits

- Ensuring k-anonymity prevents database linkage attacks with a bound of k
- Linking the datasets to external sources will not yield less than k individuals

l-Diversity

Let a q^* -block in anonymized database D*

The block is l-diverse if it contains at least l well-represented values of the sensitive attribute S

Example→

Block 1, with $q^* = [130^{**}, >50]$

has l-diversity of 1

Block 2, with $q^* = [140^{**}, <50]$

has l-diversity of 2

Block 3, with $q^* = [130^{**}, <50]$

has l-diversity of 1

Benefits

- Large values of diversity will ensure that the homogeneity and background knowledge attacks do not work
- Following the l-diversity principle you try to ensure: $\beta (q, s, D^*) \approx \alpha (q, s)$

Entropy l-diversity

A database is entropy l-diverse if for every q^* -block it holds that:

$$\sum_{s \in S} p_{(q^*, s)} \log_2 p_{(q^*, s)} \geq \log_2(l)$$

Security Protocols

MITM attack pattern (relay + substitution)

- Intercept A→B handshake.
- Replace public keys/certificates or swap nonces.


Replay

previously observed signed values when freshness is not checked.

- If either party doesn't validate a certificate chain or check that the key belongs to the claimed identity, attacker succeeds.

TLS Protocol: Handshake protocol (public-key operations)+ Record protocol (symmetric)

Handshake protocol



► A full handshake starts with ClientHello where the Client submits its connection and security parameters to the Server

- The Server replies with ServerHello that selected the parameters
- The Server sends a signed Certificate for authentication and the Client verifies the signature

► The Client and Server use public key cryptography to exchange keys. The Server sends the needed information via ServerKeyExchange and notifies that is done.

► The Client sends the needed information via ClientKeyExchange, derives the premaster secret and the master secret

► Now the Client switches encryption on and notifies the Server via ChangeCipherSpec

To verify the integrity of the handshake so far, the Client computes a MAC(handshake messages) and sends it to the Server via Finished

► The Server derives the premaster secret and the master secret, switches encryption on and notifies the Client via ChangeCipherSpec

► To verify the integrity of the handshake, the Server computes a MAC(handshake messages) and sends it to the Client via Finished

► A TLS session between Client-Server may now begin using a symmetric algorithm for bulk encryption of application data

Record Protocol

► Record protocol transports and encrypts all connection messages

- The record keeps also a unique 64-bit sequence number to avoid replays

