

Timing Attacks

Kostas Papagiannopoulos
University of Amsterdam
k.papagiannopoulos@uva.nl

Contents

Introduction

AES Tbox Implementation

Cache Timing Attacks on AES

Introduction

Introduction



- ▶ Cryptographic algorithms like AES, PRESENT and RSA are a black box, parameterized with key K , that turns plaintext P to ciphertext C

Introduction




- ▶ Cryptographic algorithms like AES, PRESENT and RSA are a black box, parameterized with key K , that turns plaintext P to ciphertext C
- ▶ Analyzing the cipher's security in the **blackbox scenario** relates to classical cryptanalysis techniques like linear and differential cryptanalysis

Introduction




- ▶ Cryptographic algorithms like AES, PRESENT and RSA are a black box, parameterized with key K , that turns plaintext P to ciphertext C
- ▶ Analyzing the cipher's security in the **blackbox scenario** relates to classical cryptanalysis techniques like linear and differential cryptanalysis
- ▶ “Can you recover the secret key by observing plaintext/ciphertext pairs?”

Introduction




- ▶ The cryptographic algorithm is typically implemented on a device such as a desktop/laptop/cellphone processor, microcontroller, FPGA, ASIC etc.

Introduction




- ▶ The cryptographic algorithm is typically implemented on a device such as a desktop/laptop/cellphone processor, microcontroller, FPGA, ASIC etc.
- ▶ We can observe certain **physical quantities** in the device's vicinity such as the cipher's execution time, the power consumption of the device and others

Introduction




- ▶ The cryptographic algorithm is typically implemented on a device such as a desktop/laptop/cellphone processor, microcontroller, FPGA, ASIC etc.
- ▶ We can observe certain **physical quantities** in the device's vicinity such as the cipher's execution time, the power consumption of the device and others
- ▶ Although they seem harmless, these quantities can provide substantial information during cryptanalysis
- ▶ In fact, they grant us limited access to the internal computations carried out by the cipher

Introduction




- ▶ We refer to this case as the **greybox scenario** and to the relevant physical quantities as **side-channels**

Introduction




- ▶ We refer to this case as the **greybox scenario** and to the relevant physical quantities as **side-channels**
- ▶ The **side-channel leakage** is any unintentional signal that offers us a blurry view of the cipher's internal computations
 - e.g. the padding error message during server communication can be considered a 1-bit side-channel leakage

Introduction



- ▶ "Can you derive the secret key by observing plaintext/ciphertext pairs **and a side-channel?**"

Introduction



- ▶ "Can you derive the secret key by observing plaintext/ciphertext pairs **and a side-channel?**"
- ▶ Secure algorithms under the blackbox scenario may not be secure under the greybox scenario
 - e.g. AES and DES are resistant to differential cryptanalysis but can be broken easily with a side-channel attack

Introduction

- ▶ We describe a simple timing attack on PIN code verification
- ▶ Assume that the following code was written (hastily) to check if a PIN matches the correct value 5902

Input: A 4-digit PIN code

Output: PIN accept or reject

```
1 Process CheckPIN(pin[4])
2 int pin_ok = 0;
3 if pin[0]==5 then
4     if pin[1]==9 then
5         if pin[2]==0 then
6             if pin[3]==2 then
7                 pin_ok = 1 ;
8             end
9         end
10    end
11 end
12 return pin_ok ;
13 EndProcess
```

Introduction

- ▶ We describe a simple timing attack on PIN code verification
- ▶ Assume that the following code was written (hastily) to check if a PIN matches the correct value 5902

Input: A 4-digit PIN code

Output: PIN accept or reject

```
1 Process CheckPIN(pin[4])
2 int pin_ok = 0;
3 if pin[0]==5 then
4     if pin[1]==9 then
5         if pin[2]==0 then
6             if pin[3]==2 then
7                 |   pin_ok = 1 ;
8             end
9         end
10    end
11 end
12 return pin_ok ;
13 EndProcess
```

- ▶ What are the execution times of the process for PIN inputs [0,1,2,3] , [5,3,0,2] , [5,9,0,0] ?

Introduction

- ▶ We describe a simple timing attack on PIN code verification
- ▶ Assume that the following code was written (hastily) to check if a PIN matches the correct value 5902

Input: A 4-digit PIN code

Output: PIN accept or reject

```
1 Process CheckPIN(pin[4])
2 int pin_ok = 0;
3 if pin[0]==5 then
4     if pin[1]==9 then
5         if pin[2]==0 then
6             if pin[3]==2 then
7                 pin_ok = 1 ;
8             end
9         end
10    end
11 end
12 return pin_ok ;
13 EndProcess
```

- ▶ What are the execution times of the process for PIN inputs [0,1,2,3] , [5,3,0,2] , [5,9,0,0] ?
- ▶ Notice that the execution time increases as we get closer to the correct PIN value [5,9,0,2]

Introduction

- ▶ We describe a simple timing attack on PIN code verification
- ▶ Assume that the following code was written (hastily) to check if a PIN matches the correct value 5902

Input: A 4-digit PIN code

Output: PIN accept or reject

```
1 Process CheckPIN(pin[4])
2 int pin_ok = 0;
3 if pin[0]==5 then
4     if pin[1]==9 then
5         if pin[2]==0 then
6             if pin[3]==2 then
7                 pin_ok = 1 ;
8             end
9         end
10    end
11 end
12 return pin_ok ;
13 EndProcess
```

- ▶ What are the execution times of the process for PIN inputs [0,1,2,3] , [5,3,0,2] , [5,9,0,0] ?
- ▶ Notice that the execution time increases as we get closer to the correct PIN value [5,9,0,2]
- ▶ Timing attacks exploit the relationship between the execution time and a secret (e.g. a PIN or a cipher key)

AES Tbox Implementation

AES Tbox

AES T-box implementation

- ▶ The standard AES code implements every operation independently
i.e. we code and optimize the key addition, sbox, shift rows and mix columns separately

AES Tbox

AES T-box implementation

- ▶ The standard AES code implements every operation independently i.e. we code and optimize the key addition, sbox, shift rows and mix columns separately
- ▶ To speed up software implementations, Rijndael proposed **merging** the sbox, shift rows and mix columns operations in lookup tables (LUTs) called **T-boxes**

AES Tbox

AES T-box implementation

- ▶ The standard AES code implements every operation independently i.e. we code and optimize the key addition, sbox, shift rows and mix columns separately
- ▶ To speed up software implementations, Rijndael proposed **merging** the sbox, shift rows and mix columns operations in lookup tables (LUTs) called **T-boxes**
- ▶ The merging led to very fast implementations, especially in 32-bit architectures
- ▶ The downside was increased memory requirements, since the implementation needed 4 lookup tables of 1024 bytes each

AES Tbox

AES T-box implementation

- ▶ The standard AES code implements every operation independently i.e. we code and optimize the key addition, sbox, shift rows and mix columns separately
- ▶ To speed up software implementations, Rijndael proposed **merging** the sbox, shift rows and mix columns operations in lookup tables (LUTs) called **T-boxes**
- ▶ The merging led to very fast implementations, especially in 32-bit architectures
- ▶ The downside was increased memory requirements, since the implementation needed 4 lookup tables of 1024 bytes each
- ▶ Software libraries often use the merged AES implementation while resource-constrained hardware keep standard implementations

AES Tbox

AES 16-byte state

$$\text{state } A = \begin{bmatrix} a_{0,0} & a_{0,1} & a_{0,2} & a_{0,3} \\ a_{1,0} & a_{1,1} & a_{1,2} & a_{1,3} \\ a_{2,0} & a_{2,1} & a_{2,2} & a_{2,3} \\ a_{3,0} & a_{3,1} & a_{3,2} & a_{3,3} \end{bmatrix}$$

AES Tbox

AES 16-byte state

$$\text{state } A = \begin{bmatrix} a_{0,0} & a_{0,1} & a_{0,2} & a_{0,3} \\ a_{1,0} & a_{1,1} & a_{1,2} & a_{1,3} \\ a_{2,0} & a_{2,1} & a_{2,2} & a_{2,3} \\ a_{3,0} & a_{3,1} & a_{3,2} & a_{3,3} \end{bmatrix}$$

Sbox operation

$$b_{i,j} = S(a_{i,j}), \quad \text{for all } i, j \in \{0, 1, 2, 3\}$$

In column-vector notation:

$$\begin{bmatrix} b_{0,j} \\ b_{1,j} \\ b_{2,j} \\ b_{3,j} \end{bmatrix} = \begin{bmatrix} S(a_{0,j}) \\ S(a_{1,j}) \\ S(a_{2,j}) \\ S(a_{3,j}) \end{bmatrix}, \quad \text{for } j = 0, 1, 2, 3$$

AES Tbox

Shift Rows operation

$$\begin{bmatrix} c_{0,j} \\ c_{1,j} \\ c_{2,j} \\ c_{3,j} \end{bmatrix} = \begin{bmatrix} b_{0,j} \\ b_{1,j-C_1} \\ b_{2,j-C_2} \\ b_{3,j-C_3} \end{bmatrix}, \text{ for } j = 0, 1, 2, 3$$

For AES-128: $C_1 = 1, C_2 = 2, C_3 = 3$ and $j - C = \text{mod}(j - C, 4)$

AES Tbox

Shift Rows operation

$$\begin{bmatrix} c_{0,j} \\ c_{1,j} \\ c_{2,j} \\ c_{3,j} \end{bmatrix} = \begin{bmatrix} b_{0,j} \\ b_{1,j-C_1} \\ b_{2,j-C_2} \\ b_{3,j-C_3} \end{bmatrix}, \text{ for } j = 0, 1, 2, 3$$

For AES-128: $C_1 = 1, C_2 = 2, C_3 = 3$ and $j - C = \text{mod}(j - C, 4)$

Mix Columns operation

$$\begin{bmatrix} d_{0,j} \\ d_{1,j} \\ d_{2,j} \\ d_{3,j} \end{bmatrix} = \begin{bmatrix} 02 & 03 & 01 & 01 \\ 01 & 02 & 03 & 01 \\ 01 & 01 & 02 & 03 \\ 03 & 01 & 01 & 02 \end{bmatrix} \begin{bmatrix} c_{0,j} \\ c_{1,j} \\ c_{2,j} \\ c_{3,j} \end{bmatrix}, \text{ for } j = 0, 1, 2, 3$$

AES Tbox

Key addition operation

$$\begin{bmatrix} e_{0,j} \\ e_{1,j} \\ e_{2,j} \\ e_{3,j} \end{bmatrix} = \begin{bmatrix} d_{0,j} \\ d_{1,j} \\ d_{2,j} \\ d_{3,j} \end{bmatrix} \oplus \begin{bmatrix} k_{0,j} \\ k_{1,j} \\ k_{2,j} \\ k_{3,j} \end{bmatrix}, \text{ for } j = 0, 1, 2, 3$$

AES Tbox

- ▶ We now combine the previous expressions backwards (for $j = 0, 1, 2, 3$)

$$\begin{bmatrix} e_{0,j} \\ e_{1,j} \\ e_{2,j} \\ e_{3,j} \end{bmatrix} = \begin{bmatrix} d_{0,j} \\ d_{1,j} \\ d_{2,j} \\ d_{3,j} \end{bmatrix} \oplus \begin{bmatrix} k_{0,j} \\ k_{1,j} \\ k_{2,j} \\ k_{3,j} \end{bmatrix} =$$

AES Tbox

- We now combine the previous expressions backwards (for $j = 0, 1, 2, 3$)

$$\begin{bmatrix} e_{0,j} \\ e_{1,j} \\ e_{2,j} \\ e_{3,j} \end{bmatrix} = \begin{bmatrix} d_{0,j} \\ d_{1,j} \\ d_{2,j} \\ d_{3,j} \end{bmatrix} \oplus \begin{bmatrix} k_{0,j} \\ k_{1,j} \\ k_{2,j} \\ k_{3,j} \end{bmatrix} = \begin{bmatrix} 02 & 03 & 01 & 01 \\ 01 & 02 & 03 & 01 \\ 01 & 01 & 02 & 03 \\ 03 & 01 & 01 & 02 \end{bmatrix} \begin{bmatrix} c_{0,j} \\ c_{1,j} \\ c_{2,j} \\ c_{3,j} \end{bmatrix} \oplus \begin{bmatrix} k_{0,j} \\ k_{1,j} \\ k_{2,j} \\ k_{3,j} \end{bmatrix} =$$

AES Tbox

- We now combine the previous expressions backwards (for $j = 0, 1, 2, 3$)

$$\begin{bmatrix} e_{0,j} \\ e_{1,j} \\ e_{2,j} \\ e_{3,j} \end{bmatrix} = \begin{bmatrix} d_{0,j} \\ d_{1,j} \\ d_{2,j} \\ d_{3,j} \end{bmatrix} \oplus \begin{bmatrix} k_{0,j} \\ k_{1,j} \\ k_{2,j} \\ k_{3,j} \end{bmatrix} = \begin{bmatrix} 02 & 03 & 01 & 01 \\ 01 & 02 & 03 & 01 \\ 01 & 01 & 02 & 03 \\ 03 & 01 & 01 & 02 \end{bmatrix} \begin{bmatrix} c_{0,j} \\ c_{1,j} \\ c_{2,j} \\ c_{3,j} \end{bmatrix} \oplus \begin{bmatrix} k_{0,j} \\ k_{1,j} \\ k_{2,j} \\ k_{3,j} \end{bmatrix} =$$

$$\begin{bmatrix} 02 & 03 & 01 & 01 \\ 01 & 02 & 03 & 01 \\ 01 & 01 & 02 & 03 \\ 03 & 01 & 01 & 02 \end{bmatrix} \begin{bmatrix} b_{0,j} \\ b_{1,j-C_1} \\ b_{2,j-C_2} \\ b_{3,j-C_3} \end{bmatrix} \oplus \begin{bmatrix} k_{0,j} \\ k_{1,j} \\ k_{2,j} \\ k_{3,j} \end{bmatrix} =$$

AES Tbox

- We now combine the previous expressions backwards (for $j = 0, 1, 2, 3$)

$$\begin{bmatrix} e_{0,j} \\ e_{1,j} \\ e_{2,j} \\ e_{3,j} \end{bmatrix} = \begin{bmatrix} d_{0,j} \\ d_{1,j} \\ d_{2,j} \\ d_{3,j} \end{bmatrix} \oplus \begin{bmatrix} k_{0,j} \\ k_{1,j} \\ k_{2,j} \\ k_{3,j} \end{bmatrix} = \begin{bmatrix} 02 & 03 & 01 & 01 \\ 01 & 02 & 03 & 01 \\ 01 & 01 & 02 & 03 \\ 03 & 01 & 01 & 02 \end{bmatrix} \begin{bmatrix} c_{0,j} \\ c_{1,j} \\ c_{2,j} \\ c_{3,j} \end{bmatrix} \oplus \begin{bmatrix} k_{0,j} \\ k_{1,j} \\ k_{2,j} \\ k_{3,j} \end{bmatrix} =$$

$$\begin{bmatrix} 02 & 03 & 01 & 01 \\ 01 & 02 & 03 & 01 \\ 01 & 01 & 02 & 03 \\ 03 & 01 & 01 & 02 \end{bmatrix} \begin{bmatrix} b_{0,j} \\ b_{1,j-C_1} \\ b_{2,j-C_2} \\ b_{3,j-C_3} \end{bmatrix} \oplus \begin{bmatrix} k_{0,j} \\ k_{1,j} \\ k_{2,j} \\ k_{3,j} \end{bmatrix} =$$

$$\begin{bmatrix} 02 & 03 & 01 & 01 \\ 01 & 02 & 03 & 01 \\ 01 & 01 & 02 & 03 \\ 03 & 01 & 01 & 02 \end{bmatrix} \begin{bmatrix} S(a_{0,j}) \\ S(a_{1,j-C_1}) \\ S(a_{2,j-C_2}) \\ S(a_{3,j-C_3}) \end{bmatrix} \oplus \begin{bmatrix} k_{0,j} \\ k_{1,j} \\ k_{2,j} \\ k_{3,j} \end{bmatrix}$$

AES Tbox

- Using the matrix multiplication rule we can re-write the expression as follows

$$\begin{bmatrix} e_{0,j} \\ e_{1,j} \\ e_{2,j} \\ e_{3,j} \end{bmatrix} = S(a_{0,j}) \begin{bmatrix} 02 \\ 01 \\ 01 \\ 03 \end{bmatrix} \oplus S(a_{1,j-c_1}) \begin{bmatrix} 03 \\ 02 \\ 01 \\ 01 \end{bmatrix} \oplus S(a_{2,j-c_2}) \begin{bmatrix} 01 \\ 03 \\ 02 \\ 01 \end{bmatrix} \oplus S(a_{3,j-c_3}) \begin{bmatrix} 01 \\ 01 \\ 03 \\ 02 \end{bmatrix} \oplus \begin{bmatrix} k_{0,j} \\ k_{1,j} \\ k_{2,j} \\ k_{3,j} \end{bmatrix}$$

AES Tbox

- ▶ Using the matrix multiplication rule we can re-write the expression as follows

$$\begin{bmatrix} e_{0,j} \\ e_{1,j} \\ e_{2,j} \\ e_{3,j} \end{bmatrix} = S(a_{0,j}) \begin{bmatrix} 02 \\ 01 \\ 01 \\ 03 \end{bmatrix} \oplus S(a_{1,j-c_1}) \begin{bmatrix} 03 \\ 02 \\ 01 \\ 01 \end{bmatrix} \oplus S(a_{2,j-c_2}) \begin{bmatrix} 01 \\ 03 \\ 02 \\ 01 \end{bmatrix} \oplus S(a_{3,j-c_3}) \begin{bmatrix} 01 \\ 01 \\ 03 \\ 02 \end{bmatrix} \oplus \begin{bmatrix} k_{0,j} \\ k_{1,j} \\ k_{2,j} \\ k_{3,j} \end{bmatrix}$$

- ▶ Observe that we can now view the column j of the AES round output $[e_{0,j} \ e_{1,j} \ e_{2,j} \ e_{3,j}]^T$ as the sum of 4 vectors with the key vector

AES Tbox

- We now define 4 lookup tables aka the **T-boxes** $T_0(\cdot)$, $T_1(\cdot)$, $T_2(\cdot)$, $T_3(\cdot)$

$$T_0(a) = \begin{bmatrix} S(a) \cdot 02 \\ S(a) \\ S(a) \\ S(a) \cdot 03 \end{bmatrix}, \quad T_1(a) = \begin{bmatrix} S(a) \cdot 03 \\ S(a) \cdot 02 \\ S(a) \\ S(a) \end{bmatrix}$$

$$T_2(a) = \begin{bmatrix} S(a) \\ S(a) \cdot 03 \\ S(a) \cdot 02 \\ S(a) \end{bmatrix}, \quad T_3(a) = \begin{bmatrix} S(a) \\ S(a) \\ S(a) \cdot 03 \\ S(a) \cdot 02 \end{bmatrix}$$

AES Tbox

- ▶ We now define 4 lookup tables aka the **T-boxes** $T_0(\cdot)$, $T_1(\cdot)$, $T_2(\cdot)$, $T_3(\cdot)$

$$T_0(a) = \begin{bmatrix} S(a) \cdot 02 \\ S(a) \\ S(a) \\ S(a) \cdot 03 \end{bmatrix}, \quad T_1(a) = \begin{bmatrix} S(a) \cdot 03 \\ S(a) \cdot 02 \\ S(a) \\ S(a) \end{bmatrix}$$

$$T_2(a) = \begin{bmatrix} S(a) \\ S(a) \cdot 03 \\ S(a) \cdot 02 \\ S(a) \end{bmatrix}, \quad T_3(a) = \begin{bmatrix} S(a) \\ S(a) \\ S(a) \cdot 03 \\ S(a) \cdot 02 \end{bmatrix}$$

- ▶ Each T-box has 256 entries since the byte-sized input a takes values in $\{0, 1, \dots, 255\}$
- ▶ Every T-box entry contains 4 rows with values thus every entry contains 4 bytes

AES Tbox

- ▶ We now define 4 lookup tables aka the **T-boxes** $T_0(\cdot)$, $T_1(\cdot)$, $T_2(\cdot)$, $T_3(\cdot)$

$$T_0(a) = \begin{bmatrix} S(a) \cdot 02 \\ S(a) \\ S(a) \\ S(a) \cdot 03 \end{bmatrix}, \quad T_1(a) = \begin{bmatrix} S(a) \cdot 03 \\ S(a) \cdot 02 \\ S(a) \\ S(a) \end{bmatrix}$$

$$T_2(a) = \begin{bmatrix} S(a) \\ S(a) \cdot 03 \\ S(a) \cdot 02 \\ S(a) \end{bmatrix}, \quad T_3(a) = \begin{bmatrix} S(a) \\ S(a) \\ S(a) \cdot 03 \\ S(a) \cdot 02 \end{bmatrix}$$

- ▶ Each T-box has 256 entries since the byte-sized input a takes values in $\{0, 1, \dots, 255\}$
- ▶ Every T-box entry contains 4 rows with values thus every entry contains 4 bytes
- ▶ In total, every T-box has a size of $256 \times 4 = 1024$ bytes. Combined, all 4 T-boxes require 4Kbytes of memory space.

AES Tbox

- ▶ Using the 4 T-boxes we can now express all the AES round operations as follows

$$\begin{bmatrix} e_{0,j} \\ e_{1,j} \\ e_{2,j} \\ e_{3,j} \end{bmatrix} = T_0(a_{0,j}) \oplus T_1(a_{1,j} - c_1) \oplus T_2(a_{2,j} - c_2) \oplus T_3(a_{3,j} - c_3) \oplus \begin{bmatrix} k_{0,j} \\ k_{1,j} \\ k_{2,j} \\ k_{3,j} \end{bmatrix}$$

AES Tbox

- ▶ Using the 4 T-boxes we can now express all the AES round operations as follows

$$\begin{bmatrix} e_{0,j} \\ e_{1,j} \\ e_{2,j} \\ e_{3,j} \end{bmatrix} = T_0(a_{0,j}) \oplus T_1(a_{1,j} - c_1) \oplus T_2(a_{2,j} - c_2) \oplus T_3(a_{3,j} - c_3) \oplus \begin{bmatrix} k_{0,j} \\ k_{1,j} \\ k_{2,j} \\ k_{3,j} \end{bmatrix}$$

- ▶ Repeating the expression above for $j = 0, 1, 2, 3$ will yield the full 16-byte cipher state
- ▶ The T-boxes have merged the sbox, shift rows and mix columns, improving speed but increasing memory requirements

Cache Timing Attacks on AES

Cache Timing Attacks on AES

Attack Idea:

- ▶ An AES implementation that uses T-boxes will perform several memory lookups

Cache Timing Attacks on AES

Attack Idea:

- ▶ An AES implementation that uses T-boxes will perform several memory lookups
- ▶ AES starts with the 16-byte plaintext P getting XORed with the round key K_{round0}

$$A = P \oplus K_{round0}$$

Cache Timing Attacks on AES

Attack Idea:

- ▶ An AES implementation that uses T-boxes will perform several memory lookups
- ▶ AES starts with the 16-byte plaintext P getting XORed with the round key K_{round0}

$$A = P \oplus K_{round0}$$

- ▶ Then the state A will be used to perform T-box lookups in memory

$$T_0(a_{0,j}), \quad T_1(a_{1,j-C_1}), \quad T_2(a_{2,j-C_2}), \quad T_3(a_{3,j-C_3}), \quad j = 0, 1, 2, 3$$

Cache Timing Attacks on AES

Attack Idea:

- ▶ An AES implementation that uses T-boxes will perform several memory lookups
- ▶ AES starts with the 16-byte plaintext P getting XORed with the round key K_{round0}

$$A = P \oplus K_{round0}$$

- ▶ Then the state A will be used to perform T-box lookups in memory

$$T_0(a_{0,j}), \quad T_1(a_{1,j-C_1}), \quad T_2(a_{2,j-C_2}), \quad T_3(a_{3,j-C_3}), \quad j = 0, 1, 2, 3$$

- ▶ Let's focus on the 1st T-box lookup for $j = 0$

$$T_0(a_{0,0}) = T_0(p_{0,0} \oplus k_{0,0})$$

Cache Timing Attacks on AES

Attack Idea:

- ▶ An AES implementation that uses T-boxes will perform several memory lookups
- ▶ AES starts with the 16-byte plaintext P getting XORed with the round key K_{round0}

$$A = P \oplus K_{round0}$$

- ▶ Then the state A will be used to perform T-box lookups in memory

$$T_0(a_{0,j}), \quad T_1(a_{1,j-C_1}), \quad T_2(a_{2,j-C_2}), \quad T_3(a_{3,j-C_3}), \quad j = 0, 1, 2, 3$$

- ▶ Let's focus on the 1st T-box lookup for $j = 0$

$$T_0(a_{0,0}) = T_0(p_{0,0} \oplus k_{0,0})$$

- ▶ Notice that if there is any relationship between the execution time of this operation and the lookup index $a_{0,0}$ then measuring time reveals the value of index $a_{0,0}$

Cache Timing Attacks on AES

Attack Idea:

- ▶ An AES implementation that uses T-boxes will perform several memory lookups
- ▶ AES starts with the 16-byte plaintext P getting XORed with the round key K_{round0}

$$A = P \oplus K_{round0}$$

- ▶ Then the state A will be used to perform T-box lookups in memory

$$T_0(a_{0,j}), \quad T_1(a_{1,j-C_1}), \quad T_2(a_{2,j-C_2}), \quad T_3(a_{3,j-C_3}), \quad j = 0, 1, 2, 3$$


- ▶ Let's focus on the 1st T-box lookup for $j = 0$

$$T_0(a_{0,0}) = T_0(p_{0,0} \oplus k_{0,0})$$

- ▶ Notice that if there is any relationship between the execution time of this operation and the lookup index $a_{0,0}$ then measuring time reveals the value of index $a_{0,0}$
- ▶ Any information about $a_{0,0}$ leads to information about $k_{0,0} = a_{0,0} \oplus p_{0,0}$ since $p_{0,0}$ is publicly known


Cache Timing Attacks on AES

- ▶ Processors utilize data and instruction caches to speedup memory access
- ▶ Most systems have a memory hierarchy to ensure spatial and temporal locality




Cache Timing Attacks on AES

- ▶ To perform e.g. the AES T-box lookup $T_0(0)$ we need to transfer a specific memory element of T_0 based on index $a_{0,0} = 0$ to the registers
- ▶ If $T_0(0)$ is not stored in the cache then we need to fetch it from the main memory




Cache Timing Attacks on AES

- ▶ To achieve good spatial locality we will (probably) transfer also several continuous values of table T_0 to the cache
- ▶ This type of data access is a **cache miss** and it takes a large amount of time



Cache Timing Attacks on AES

- ▶ Some subsequent memory accesses will be faster
- ▶ For example accessing $T_0(7)$ is a **cache hit** and it takes a small amount of time
- ▶ Thus there is a dependency between the index $a_{0,0}$ and the execution time
- ▶ Remember that $a_{0,0}$ relates to the key $k_{0,0}$



Cache Attacks on AES

This timing attack works in two phases: profile phase and attack phase

Profile phase:

1. Setup a server that can encrypt with AES and uses a known key k

Cache Attacks on AES

This timing attack works in two phases: profile phase and attack phase

Profile phase:

1. Setup a server that can encrypt with AES and uses a known key k
2. Send a large number (n) of random plaintexts to the server, encrypt them and measure the AES execution times t_1, t_2, \dots, t_n

Cache Attacks on AES

This timing attack works in two phases: profile phase and attack phase

Profile phase:

1. Setup a server that can encrypt with AES and uses a known key k
2. Send a large number (n) of random plaintexts to the server, encrypt them and measure the AES execution times t_1, t_2, \dots, t_n
3. Estimate the average execution time μ

$$\mu = \frac{1}{n} \sum_{i=1}^n t_i$$

Cache Attacks on AES

This timing attack works in two phases: profile phase and attack phase

Profile phase:

1. Setup a server that can encrypt with AES and uses a known key k
2. Send a large number (n) of random plaintexts to the server, encrypt them and measure the AES execution times t_1, t_2, \dots, t_n
3. Estimate the average execution time μ

$$\mu = \frac{1}{n} \sum_{i=1}^n t_i$$

4. Focus on a specific plaintext byte, say $p_{0,0}$ and group the time measurements according to the value of $p_{0,0}$

$$G_0 = \{t_i \text{ s.t. } p_{0,0} = 0\}, \quad G_1 = \{t_i \text{ s.t. } p_{0,0} = 1\}, \dots, \quad G_{255} = \{t_i \text{ s.t. } p_{0,0} = 255\}$$

Cache Timing Attacks on AES

5. Compute the average of every group

$$\mu_0 = \frac{1}{|G_0|} \sum_{t_i \in G_0} t_i, \quad \mu_1 = \frac{1}{|G_1|} \sum_{t_i \in G_1} t_i, \quad \dots, \quad \mu_{255} = \frac{1}{|G_{255}|} \sum_{t_i \in G_{255}} t_i$$

Cache Timing Attacks on AES

5. Compute the average of every group

$$\mu_0 = \frac{1}{|G_0|} \sum_{t_i \in G_0} t_i, \quad \mu_1 = \frac{1}{|G_1|} \sum_{t_i \in G_1} t_i, \quad \dots, \quad \mu_{255} = \frac{1}{|G_{255}|} \sum_{t_i \in G_{255}} t_i$$

6. Compute the difference between group average $\mu_{p_{0,0}}$ and the total average μ

$$\delta_0 = \mu_0 - \mu, \quad \delta_1 = \mu_1 - \mu, \quad \dots, \quad \delta_{255} = \mu_{255} - \mu$$

Cache Timing Attacks on AES

5. Compute the average of every group

$$\mu_0 = \frac{1}{|G_0|} \sum_{t_i \in G_0} t_i, \quad \mu_1 = \frac{1}{|G_1|} \sum_{t_i \in G_1} t_i, \quad \dots, \quad \mu_{255} = \frac{1}{|G_{255}|} \sum_{t_i \in G_{255}} t_i$$

6. Compute the difference between group average $\mu_{p_{0,0}}$ and the total average μ

$$\delta_0 = \mu_0 - \mu, \quad \delta_1 = \mu_1 - \mu, \quad \dots, \quad \delta_{255} = \mu_{255} - \mu$$

7. Find the maximum difference δ_{max} and the value p_{max} that caused it

$$p_{max} = \text{argmax}(\delta_0, \delta_1, \dots, \delta_{255})$$

Thus when the server receives plaintext input $p_{0,0} = p_{max}$, the lookup performed is (on average) the slowest

Cache Timing Attacks on AES

5. Compute the average of every group

$$\mu_0 = \frac{1}{|G_0|} \sum_{t_i \in G_0} t_i, \quad \mu_1 = \frac{1}{|G_1|} \sum_{t_i \in G_1} t_i, \quad \dots, \quad \mu_{255} = \frac{1}{|G_{255}|} \sum_{t_i \in G_{255}} t_i$$

6. Compute the difference between group average $\mu_{p_{0,0}}$ and the total average μ

$$\delta_0 = \mu_0 - \mu, \quad \delta_1 = \mu_1 - \mu, \quad \dots, \quad \delta_{255} = \mu_{255} - \mu$$

7. Find the maximum difference δ_{max} and the value p_{max} that caused it

$$p_{max} = \text{argmax}(\delta_0, \delta_1, \dots, \delta_{255})$$


Thus when the server receives plaintext input $p_{0,0} = p_{max}$, the lookup performed is (on average) the slowest

8. During the profiling phase we know the server key $k_{0,0}$ and we can use it and find which is the exact lookup index a_{max} that caused the slowest lookup

$$a_{0,0} = k_{0,0} \oplus p_{0,0} \iff a_{max} = k_{0,0} \oplus p_{max}$$


Cache Timing Attacks on AES

- We plot the differences δ for every possible value of the plaintext input byte $p_{0,0}$



Cache Timing Attacks on AES

- ▶ We plot the differences δ for every possible value of the plaintext input byte $p_{0,0}$



- ▶ The timing attack assumes that the lookup index a_{max} that caused the slowest lookup will be the same in other servers with the same hardware and AES implementation

Cache Attacks on AES

Attack phase:

1. Attack a server that uses the **exact same** hardware and AES implementation

Cache Attacks on AES

Attack phase:

1. Attack a server that uses the **exact same** hardware and AES implementation
2. Naturally, the server key k' is unknown

Cache Attacks on AES

Attack phase:

1. Attack a server that uses the **exact same** hardware and AES implementation
2. Naturally, the server key k' is unknown
3. Send a large number (n') of random plaintexts to the server, encrypt them and measure the AES execution times $t'_1, t'_2, \dots, t'_{n'}$

Cache Attacks on AES

Attack phase:

1. Attack a server that uses the **exact same** hardware and AES implementation
2. Naturally, the server key k' is unknown
3. Send a large number (n') of random plaintexts to the server, encrypt them and measure the AES execution times $t'_1, t'_2, \dots, t'_{n'}$
4. Estimate the average execution time μ'

$$\mu' = \frac{1}{n'} \sum_{i=1}^{n'} t'_i$$

Cache Attacks on AES

Attack phase:

1. Attack a server that uses the **exact same** hardware and AES implementation
2. Naturally, the server key k' is unknown
3. Send a large number (n') of random plaintexts to the server, encrypt them and measure the AES execution times $t'_1, t'_2, \dots, t'_{n'}$
4. Estimate the average execution time μ'

$$\mu' = \frac{1}{n'} \sum_{i=1}^{n'} t'_i$$

5. Focus on a specific plaintext byte, say $p'_{0,0}$ and group the time measurements according to the value of $p'_{0,0}$

$$G'_0 = \{t'_i \text{ s.t. } p'_{0,0} = 0\}, \quad G'_1 = \{t'_i \text{ s.t. } p'_{0,0} = 1\}, \dots, \quad G'_{255} = \{t'_i \text{ s.t. } p'_{0,0} = 255\}$$

Cache Timing Attacks on AES

5. Compute the average of every group

$$\mu'_0 = \frac{1}{|G'_0|} \sum_{t'_i \in G'_0} t'_i, \quad \mu'_1 = \frac{1}{|G'_1|} \sum_{t'_i \in G'_1} t'_i, \quad \dots, \quad \mu'_{255} = \frac{1}{|G'_{255}|} \sum_{t'_i \in G'_{255}} t'_i$$

Cache Timing Attacks on AES

5. Compute the average of every group

$$\mu'_0 = \frac{1}{|G'_0|} \sum_{t'_i \in G'_0} t'_i, \quad \mu'_1 = \frac{1}{|G'_1|} \sum_{t'_i \in G'_1} t'_i, \quad \dots, \quad \mu'_{255} = \frac{1}{|G'_{255}|} \sum_{t'_i \in G'_{255}} t'_i$$

6. Compute the difference between group average $\mu'_{p'_{0,0}}$ and the total average μ'

$$\delta'_0 = \mu'_0 - \mu', \quad \delta'_1 = \mu'_1 - \mu', \quad \dots, \quad \delta'_{255} = \mu'_{255} - \mu'$$

Cache Timing Attacks on AES

5. Compute the average of every group

$$\mu'_0 = \frac{1}{|G'_0|} \sum_{t'_i \in G'_0} t'_i, \quad \mu'_1 = \frac{1}{|G'_1|} \sum_{t'_i \in G'_1} t'_i, \quad \dots, \quad \mu'_{255} = \frac{1}{|G'_{255}|} \sum_{t'_i \in G'_{255}} t'_i$$

6. Compute the difference between group average $\mu'_{p'_{0,0}}$ and the total average μ'

$$\delta'_0 = \mu'_0 - \mu', \quad \delta'_1 = \mu'_1 - \mu', \quad \dots, \quad \delta'_{255} = \mu'_{255} - \mu'$$

7. Find the maximum difference δ'_{max} and the value p_{max} that caused it

$$p_{max} = \text{argmax}(\delta'_0, \delta'_1, \dots, \delta'_{255})$$

Thus when the server receives plaintext input $p'_{0,0} = p'_{max}$, the lookup performed is (on average) the slowest

Cache Timing Attacks on AES

5. Compute the average of every group

$$\mu'_0 = \frac{1}{|G'_0|} \sum_{t'_i \in G'_0} t'_i, \quad \mu'_1 = \frac{1}{|G'_1|} \sum_{t'_i \in G'_1} t'_i, \quad \dots, \quad \mu'_{255} = \frac{1}{|G'_{255}|} \sum_{t'_i \in G'_{255}} t'_i$$

6. Compute the difference between group average $\mu'_{p'_{0,0}}$ and the total average μ'

$$\delta'_0 = \mu'_0 - \mu', \quad \delta'_1 = \mu'_1 - \mu', \quad \dots, \quad \delta'_{255} = \mu'_{255} - \mu'$$

7. Find the maximum difference δ'_{max} and the value p_{max} that caused it

$$p_{max} = \text{argmax}(\delta'_0, \delta'_1, \dots, \delta'_{255})$$

Thus when the server receives plaintext input $p'_{0,0} = p'_{max}$, the lookup performed is (on average) the slowest

8. During the profiling phase we have found the lookup index a_{max} that causes the slowest lookup

Cache Timing Attacks on AES

5. Compute the average of every group

$$\mu'_0 = \frac{1}{|G'_0|} \sum_{t'_i \in G'_0} t'_i, \quad \mu'_1 = \frac{1}{|G'_1|} \sum_{t'_i \in G'_1} t'_i, \quad \dots, \quad \mu'_{255} = \frac{1}{|G'_{255}|} \sum_{t'_i \in G'_{255}} t'_i$$

6. Compute the difference between group average $\mu'_{p'_{0,0}}$ and the total average μ'

$$\delta'_0 = \mu'_0 - \mu', \quad \delta'_1 = \mu'_1 - \mu', \quad \dots, \quad \delta'_{255} = \mu'_{255} - \mu'$$

7. Find the maximum difference δ'_{max} and the value p_{max} that caused it

$$p_{max} = \text{argmax}(\delta'_0, \delta'_1, \dots, \delta'_{255})$$

Thus when the server receives plaintext input $p'_{0,0} = p'_{max}$, the lookup performed is (on average) the slowest

8. During the profiling phase we have found the lookup index a_{max} that causes the slowest lookup
9. Since the AES implementation and server hardware are the same:

$$a_{max} = a'_{max} \iff a_{max} = k'_{0,0} \oplus p'_{max} \iff k'_{0,0} = a_{max} \oplus p'_{max}$$

Cache Timing Attacks on AES

Notes on timing attacks

- ▶ Timing attacks can be executed remotely over the network
- ▶ The majority of LUT-based AES implementations can be broken. Many cryptographic libraries, including OpenSSL, were found vulnerable, regardless of the server hardware.
- ▶ Even if the server does not provide timestamps in the packets, the attacker can measure the packet round-trip-time over the network

Cache Timing Attacks on AES

Countermeasures

- ▶ Inserting random delays during the AES computations can make the attack harder but not impossible because of the Central Limit theorem
 - ▶ Let time measurements T_1, T_2, \dots, T_n , all with mean μ and standard deviation σ (i.e. the measurements are noisy)
 - ▶ Let also their average $\bar{T} = \frac{1}{n} \sum_{i=1}^n T_i$
 - ▶ Then, for large n , \bar{T} follows a normal distribution $\mathcal{N}\left(\mu, \frac{\sigma}{\sqrt{n}}\right)$

Cache Timing Attacks on AES

Countermeasures

- ▶ Inserting random delays during the AES computations can make the attack harder but not impossible because of the Central Limit theorem
 - ▶ Let time measurements T_1, T_2, \dots, T_n , all with mean μ and standard deviation σ (i.e. the measurements are noisy)
 - ▶ Let also their average $\bar{T} = \frac{1}{n} \sum_{i=1}^n T_i$
 - ▶ Then, for large n , \bar{T} follows a normal distribution $\mathcal{N}\left(\mu, \frac{\sigma}{\sqrt{n}}\right)$
- ▶ Thus, as we increase the number n of time measurements, the noise of the average measurement decreases by factor \sqrt{n} .

Cache Timing Attacks on AES

Countermeasures

- ▶ Ensure constant time execution of the encryption/decryption algorithm
 - ▶ Use custom hardware that enforces constant-time memory access – yet this hardware will be slower than average
 - ▶ Implement the whole of AES without lookup tables – yet this may damage performance, especially in software implementations
 - ▶ Fully understand cache behavior and avoid time-dependent hits and misses during memory lookups – yet this can be very challenging due to the proprietary CPU design