



**Multi-Agent Systems**

# **Game Theory**

Mina Young Pedersen

[m.y.pedersen@uva.nl](mailto:m.y.pedersen@uva.nl)

**Fall 2025**

# Decision Theory

- **Decision theory** is about deciding what is best when one agent acts against nature (not other agents)
- Example: Two lotteries, you can enter one.
  - Lottery 1: 50% chance to win 150€
  - Lottery 2: 90% chance to win 100€
  - Which one would you choose?
- You want to **maximize expected utility**

$$\mathbb{E}[u(a)] = \sum_o (u(a, o) \cdot \mathbb{P}[o])$$

utility of acting  $a$   
when the state is  $o$

probability that  $o$   
occurs

expected utility

possible state of  
nature

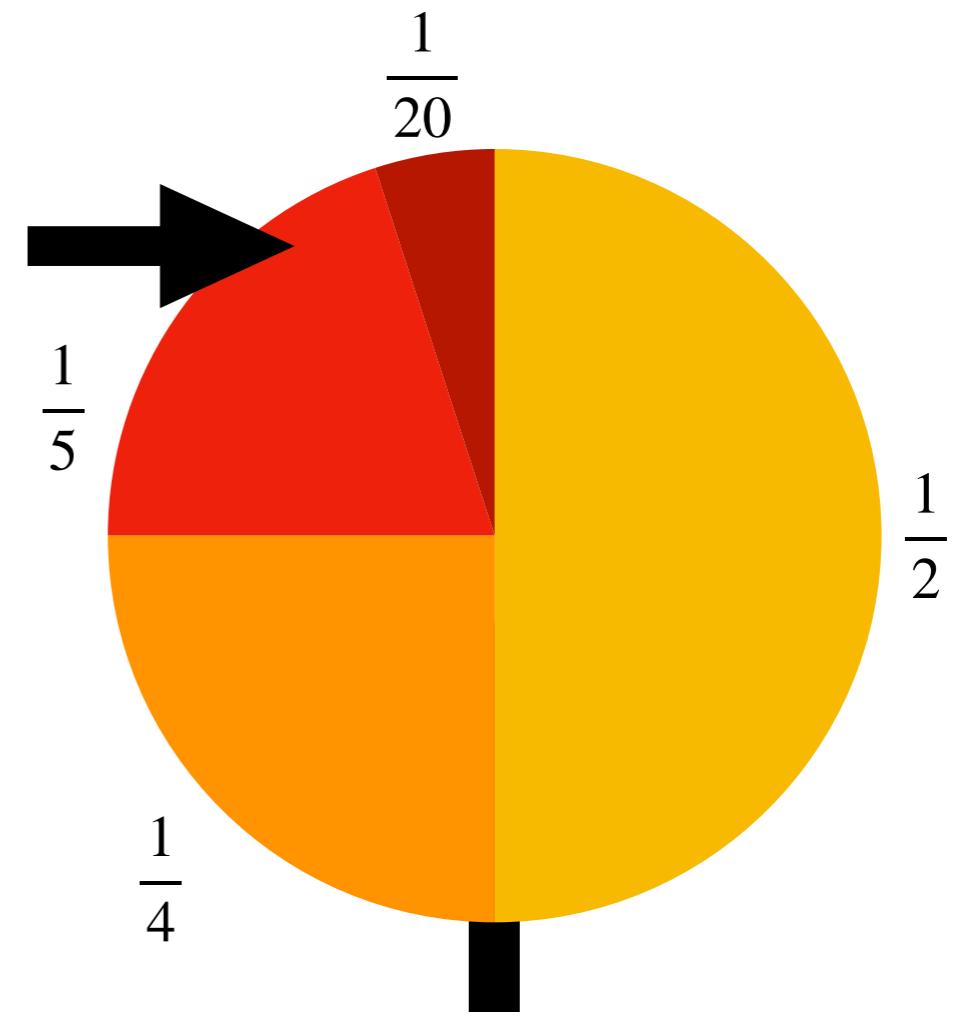


$$\mathbb{E}[u(L1)] = 150 \cdot \frac{1}{2} = 75$$

$$\mathbb{E}[u(L2)] = 100 \cdot \frac{9}{10} = 90$$

# Decision Theory: Another Example

$$\mathbb{E}[u(a)] = \sum_o (u(a, o) \cdot \mathbb{P}[o])$$



$$\mathbb{E}[u(\text{spinning the wheel})] =$$

$$u(\text{spinning, win 0}) \cdot \mathbb{P}[\text{win 0}] + u(\text{spinning, win 5}) \cdot \mathbb{P}[\text{win 5}] + \dots =$$

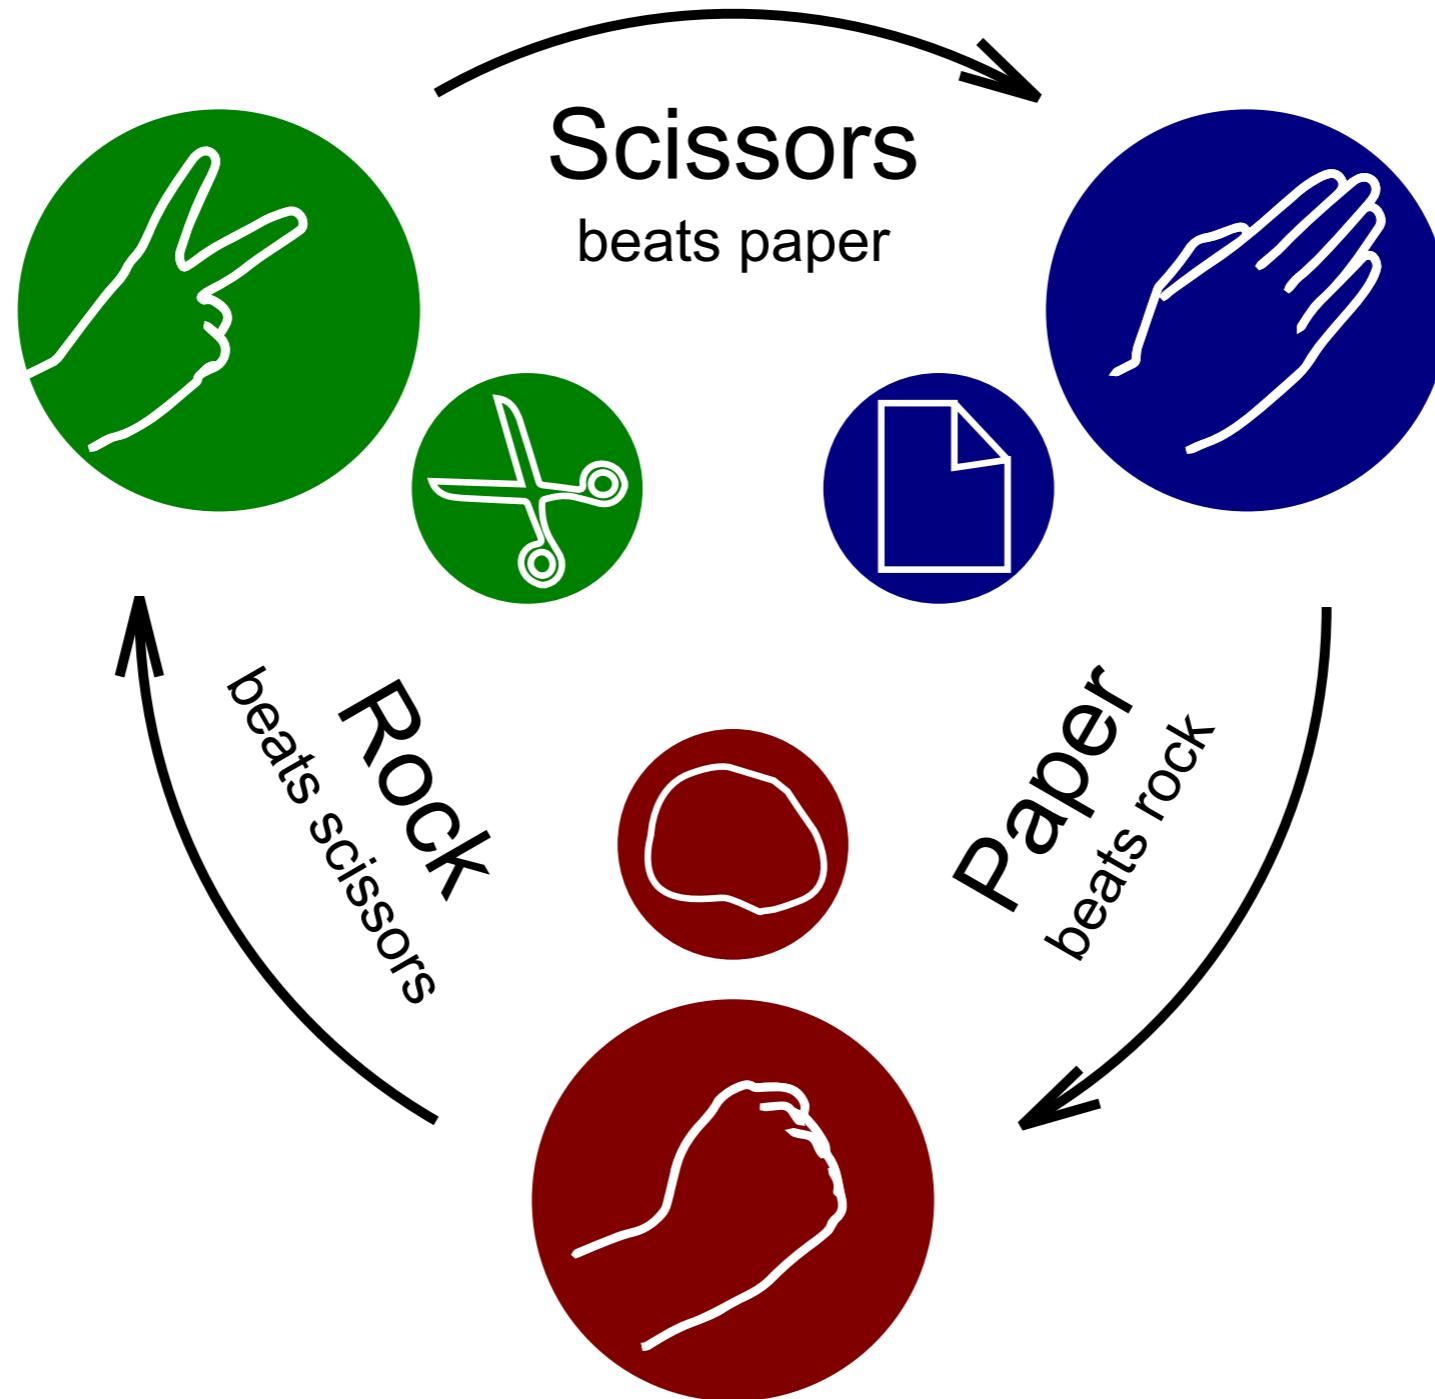
$$0 \cdot \frac{1}{2} + 5 \cdot \frac{1}{4} + 10 \cdot \frac{1}{5} + 100 \cdot \frac{1}{20} = \frac{33}{4} = 8,25$$

Coming up: Playing with others...

# What is Game Theory?

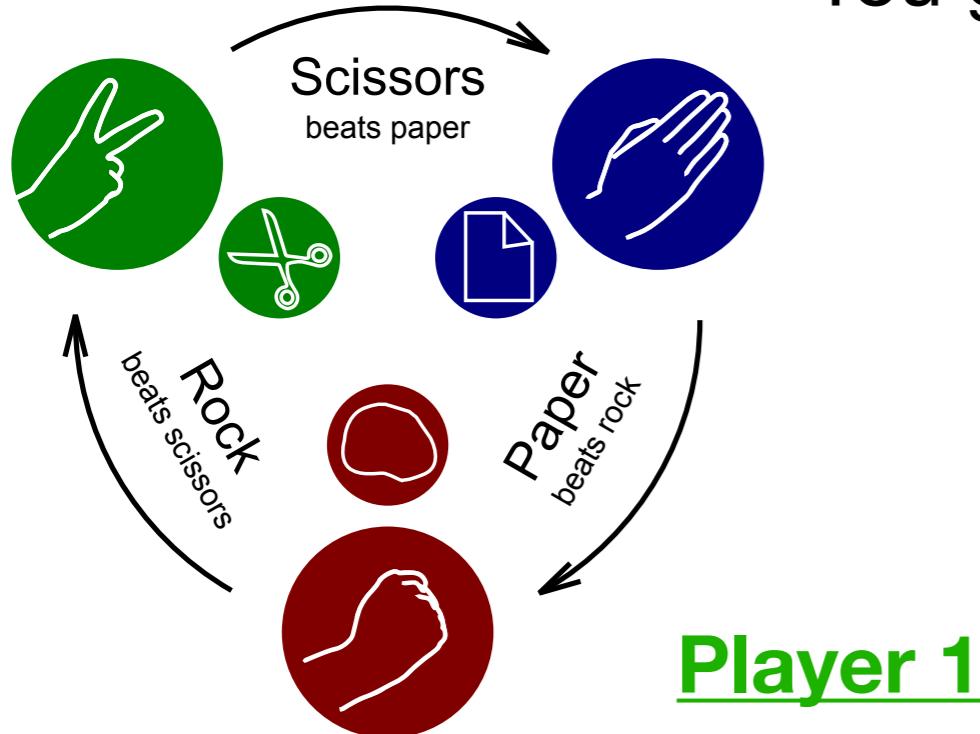
- **Game theory** is about interactions among independent, self-interested agents
- In a way: thinking about conflict, natural phenomena or daily-life scenarios as games
- **Cooperative** and **non-cooperative** game theory
  - In this course: non-cooperative
  - Players are:
    1. **Intelligent** (reason perfectly and quickly)
    2. **Rational** (always want to maximize their payoff)
    3. **Selfish** (only care for their own payoff)

# Example: Rock Paper Scissors



# Example: Rock Paper Scissors

You get 1 point for winning, and -1 for losing.



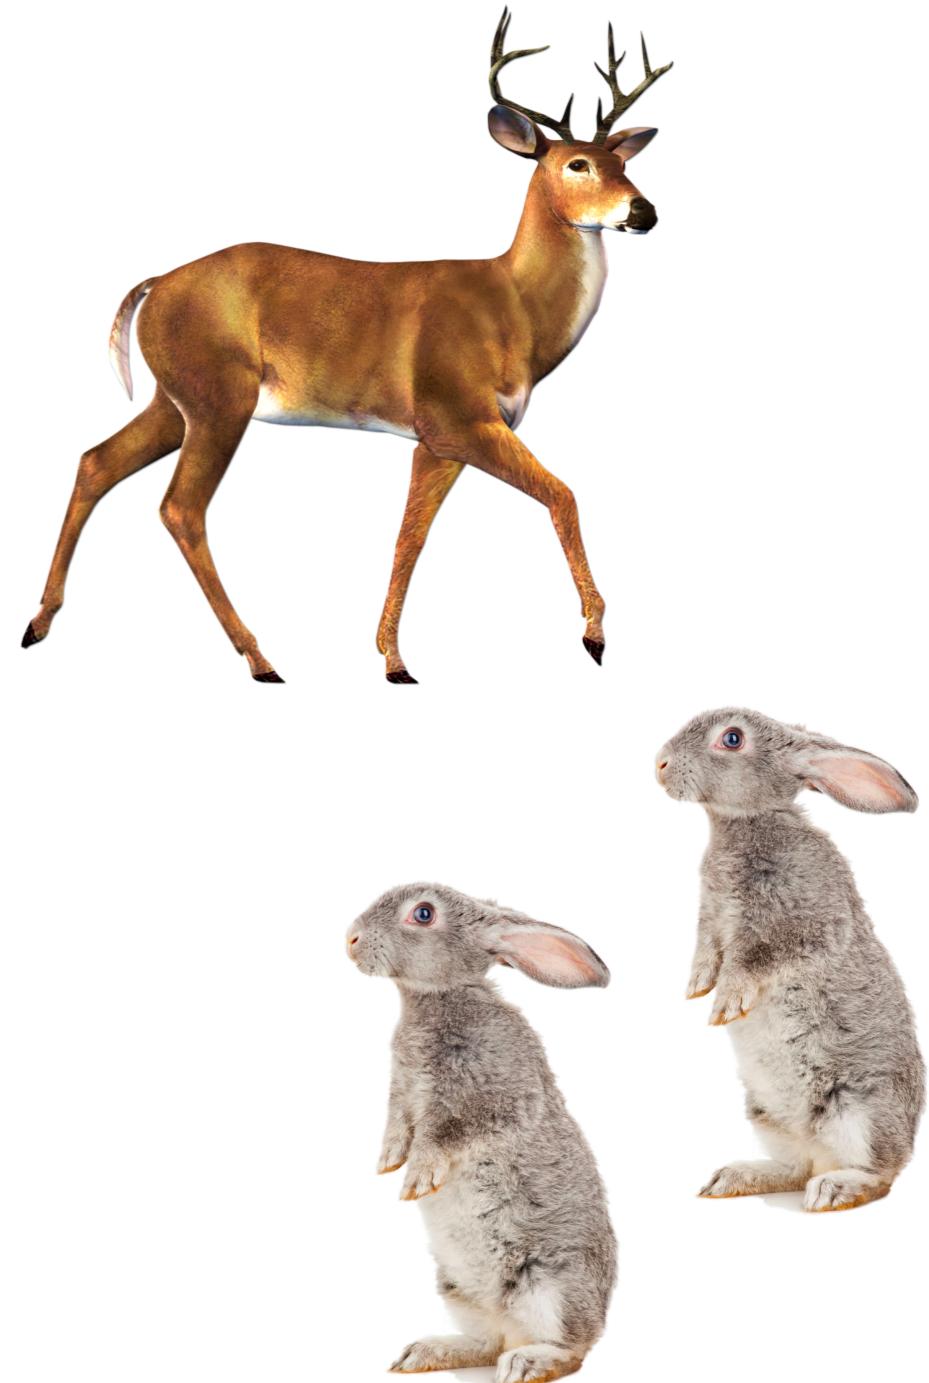
		<u>Player 2</u>		
		Rock	Paper	Scissors
<u>Player 1</u>	Rock	(0,0)	(-1,1)	(1,-1)
	Paper	(1,-1)	(0,0)	(-1,1)
	Scissors	(-1,1)	(1,-1)	(0,0)

This way of presenting a game is called a **normal form game**.

Generally, we assume Player 1 is the row player, and Player 2 is the column player.

# Another Example: Stag Hunt

- Two hunters
- Have to decide what to hunt: one stag or two hares
- A stag takes two to catch: hunting a stag solo is certain failure
- One hunter alone can catch both hares, but if both hunters go for the hares they end up with one each
- Even combined, the hares are worth less than the stag



# Another Example: Stag Hunt



Hunter 1



Hunter 2

Stag      Hare

Stag

(10,10)

(0,6)

Hare

(6,0)

(3,3)

# Games in Normal Form

- A game in normal form consists of **players** who can take **actions**, which lead to **payoffs**

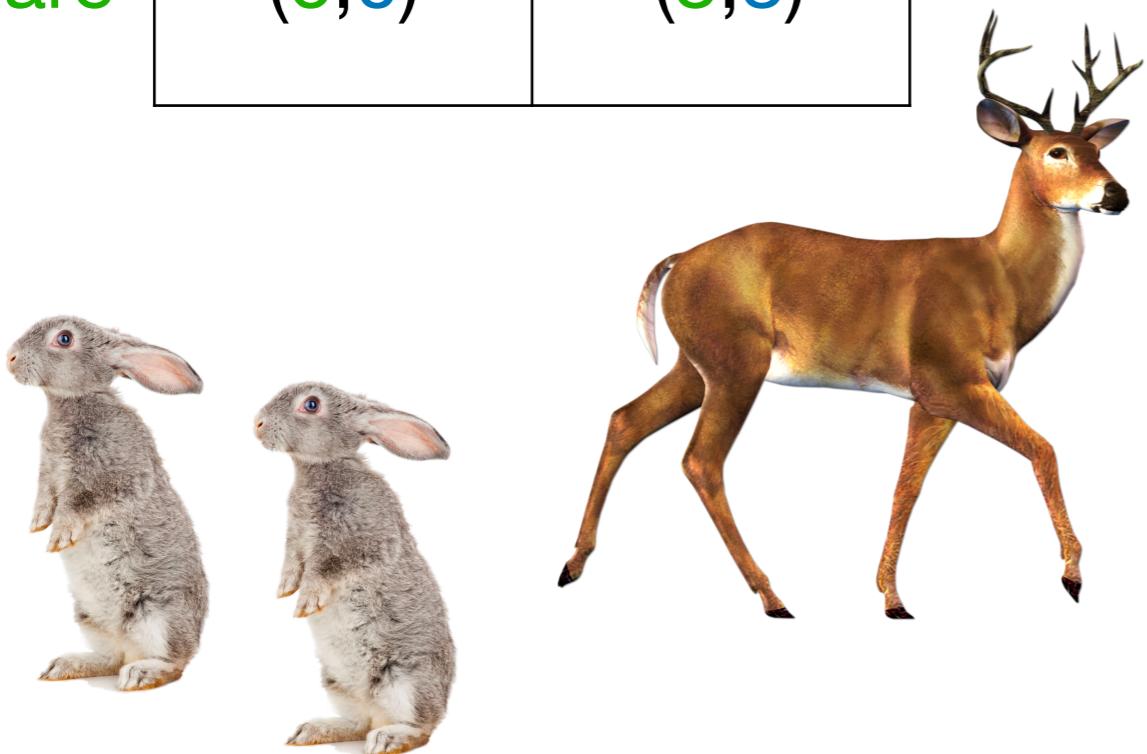
## Definition (Normal-Form Game)

A **normal-form game** is a tuple  $(N, A, u)$ , where:

- $N = \{1, \dots, n\}$  is a finite set of agents, or players
- $A_i$  is a finite set of actions available to player  $i$
- $a = (a_1, \dots, a_n)$  is an action profile
- $A = A_1 \times \dots \times A_n$  is the set of all action profiles
- $u_i : A \rightarrow \mathbb{R}$  is a utility (or payoff) function for player  $i$
- $u = (u_1, \dots, u_n)$  is a utility profile

# Back to the Stag Hunt

		<u>Hunter 2</u>	
<u>Hunter 1</u>	Stag	Hare	
Stag	(10,10)	(0,6)	
Hare	(6,0)	(3,3)	



- $N = \{1,2\}$
- $A_1 = \{Stag, Hare\}$
- $A_2 = \{Stag, Hare\}$
- $A = \{(Stag, Stag), (Stag, Hare), (Hare, Stag), (Hare, Hare)\}$
- $u_1(Stag, Stag) = 10$
- $u_2(Stag, Stag) = 10$
- $u_1(Stag, Hare) = 0$
- $u_2(Stag, Hare) = 6$
- (...)

# What Should Players Do: Strategies

- A **pure strategy** for an agent is to select a single action and play it

- $N = \{1, \dots, n\}$  is a finite set of agents, or players
- $A_i$  is a finite set of actions available to player  $i$
- $a = (a_1, \dots, a_n)$  is an action profile
- $A = A_1 \times \dots \times A_n$  is the set of all action profiles
- $u_i : A \rightarrow \mathbb{R}$  is a utility (or payoff) function for player  $i$
- $u = (u_1, \dots, u_n)$  is a utility profile

## Definition (Pure Strategy)

- $S_i = A_i$  are the pure strategies of player  $i$
- $s = (s_1, \dots, s_n)$  is a strategy profile
- $S = S_1 \times \dots \times S_n$  is the set of all strategy profiles
- $u_i(s) = u_i(a)$  is the utility of  $i$  wrt. strategy profile  $s$
- $s_{-i} = (s_1, \dots, s_{i-1}, s_{i+1}, \dots, s_n)$  is  $s$  without  $s_i$
- $s = (s_i, s_{-i})$  is another way to write  $s$

# What Should Players Do: Solution Concepts

- A **solution concept** describes what strategies we might expect the players will adopt
  - And therefore, the result of the game

# Dominance Among Strategies

- A player has a **dominated strategy** if the player would always be as well (or even better) off playing something else

# Strict Dominance

## Definition (Strict Dominance Among Strategies)

- Strategy  $s_i$  **strictly dominates**  $s'_i$  if  $u_i(s_i, s_{-i}) > u_i(s'_i, s_{-i})$ , for any profile  $s_{-i}$  of other agents' strategies
- Strategy  $s_i$  is **strictly dominant** if it strictly dominates every other strategy  $s'_i$

# Example

	L	C	R
T	(3,0)	(2,1)	(0,0)
M	(1,1)	(1,1)	(5,0)
B	(0,1)	(4,2)	(0,1)

- For Player 1, does T strictly dominate M?

- If P2 plays L:  $3 > 1$
- If P2 plays C:  $2 > 1$
- If P2 plays R:  $0 < 5$
- No

- For Player 2, does C strictly dominate L?

- If P1 plays T:  $1 > 0$
- If P1 plays M:  $1 = 1$
- If P1 plays B:  $2 > 1$
- No

# Example

	L	C	R
T	(3,0)	(2,1)	(0,0)
M	(1,1)	(1,1)	(5,0)
B	(0,1)	(4,2)	(0,1)

- For Player 2, does C strictly dominate R?

- If P1 plays T:  $1 > 0$  
- If P1 plays M:  $1 > 0$  
- If P1 plays B:  $2 > 1$  
- Yes

Player 2 should always play C over R...

...this suggests a way of filtering out strategies

# Eliminating Strictly Dominated Strategies

**Definition (Iterated Elimination of Strictly Dominated Strategies, or IESDS)**

- For each player, go through its strategies. If any are strictly dominated, eliminate them. Repeat until there are no further strategies to eliminate.

# Back to the Example

- What can we eliminate here?
  - C strictly dominates R:
    - Eliminate R

	L	C	R
T	(3,0)	(2,1)	(0,0)
M	(1,1)	(1,1)	(5,0)
B	(0,1)	(4,2)	(0,1)

# Back to the Example

- What can we eliminate here?
  - C strictly dominates R:
    - Eliminate R
  - T strictly dominates M:
    - Eliminate M

	L	C
T	$(3, 0)$	$(2, 1)$
M	$(1, 1)$	$(1, 1)$
B	$(0, 1)$	$(4, 2)$

# Back to the Example

- What can we eliminate here?
  - C strictly dominates R:
    - Eliminate R
  - T strictly dominates M:
    - Eliminate M
  - C strictly dominates L:
    - Eliminate L

	L	C
T	(3,0)	(2,1)
B	(0,1)	(4,2)

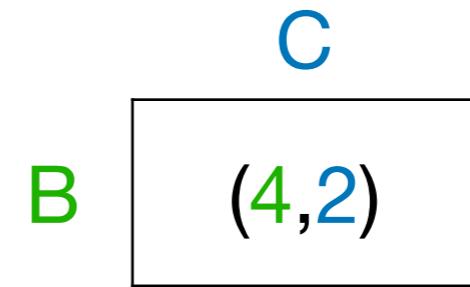
# Back to the Example

- What can we eliminate here?
  - C strictly dominates R:
    - Eliminate R
  - T strictly dominates M:
    - Eliminate M
  - C strictly dominates L:
    - Eliminate L
  - B strictly dominates T:
    - Eliminate T

	C	
T	(2, 1)	
B		(4, 2)

# Back to the Example

- What can we eliminate here?
  - C strictly dominates R:
    - Eliminate R
  - T strictly dominates M:
    - Eliminate M
  - C strictly dominates L:
    - Eliminate L
  - B strictly dominates T:
    - Eliminate T
  - Nothing left:
    - Stop



# Back to the Example

- What can we eliminate here?
  - C strictly dominates R:
    - Eliminate R
  - T strictly dominates M:
    - Eliminate M
  - C strictly dominates L:
    - Eliminate L
  - B strictly dominates T:
    - Eliminate T
  - Nothing left:
    - Stop
- IESDS solution: (B, C)

	L	C	R
T	(3,0)	(2,1)	(0,0)
M	(1,1)	(1,1)	(5,0)
B	(0,1)	(4,2)	(0,1)

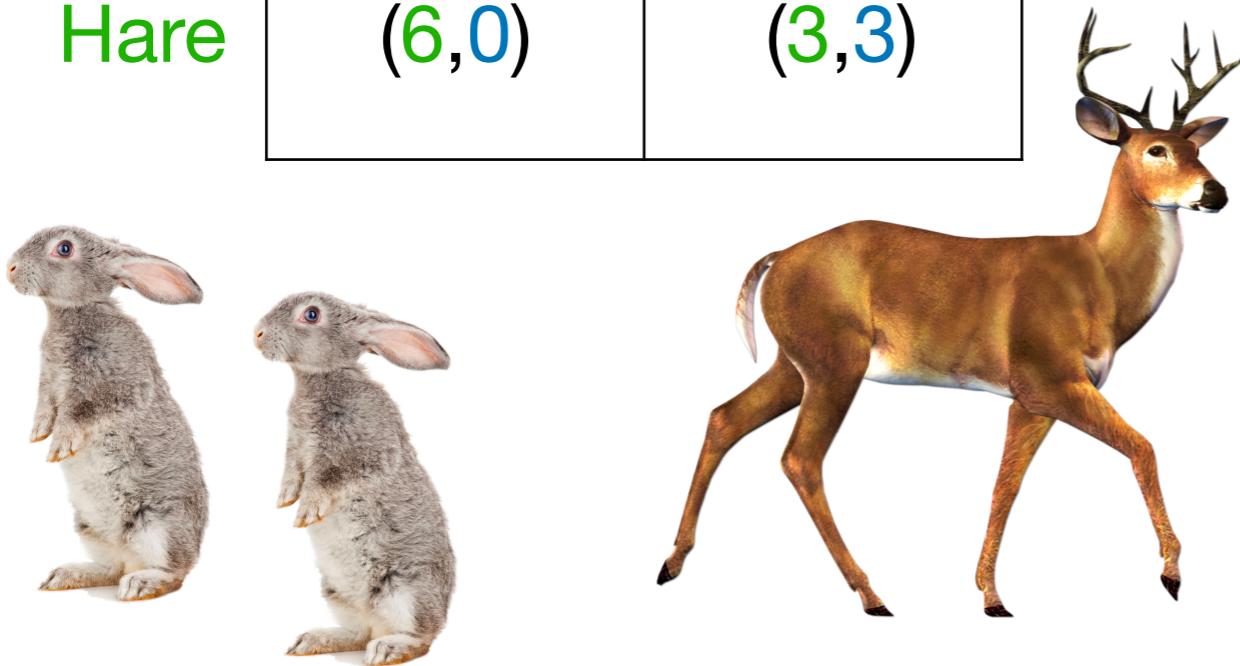
# More on IESDS

- The order in which strictly dominated strategies are eliminated does not matter
- We say that a game is **solvable** when IESDS results in a unique strategy profile

# Return to the Stag Hunt

		<u>Hunter 2</u>	
		Stag	Hare
<u>Hunter 1</u>	Stag	(10,10)	(0,6)
	Hare	(6,0)	(3,3)

- Are there any strictly dominated strategies to eliminate here?
  - No



Not all games are solvable...

# Weak Dominance

## Definition (Weak Dominance Among Strategies)

- Strategy  $s_i$  **weakly dominates**  $s'_i$  if:
  - $u_i(s_i, s_{-i}) \geq u_i(s'_i, s_{-i})$ , for any profile  $s_{-i}$  of other agents' strategies, and
  - there is some strategy profile  $\tilde{s}_{-i}$  such that  $u_i(s_i, \tilde{s}_{-i}) > u_i(s'_i, \tilde{s}_{-i})$
- Strategy  $s_i$  is **weakly dominant** if it weakly dominates every other strategy  $s'_i$

# Example

	L	C	R
T	(3,0)	(2,1)	(0,0)
M	(1,1)	(1,1)	(5,0)
B	(0,1)	(4,2)	(0,1)

- For Player 2, does C weakly dominate L?

- If P1 plays T:  $1 > 0$  
- If P1 plays M:  $1 = 1$  
- If P1 plays B:  $2 > 1$  
- Yes

# More on Weakly Dominated Strategies

- We can eliminate weakly dominated strategies, just like we eliminate strictly dominated strategies
- This time, though, the order in which we eliminate matters

This will be further explored in the homework...

# Pareto Optimality

- In a **Pareto optimal outcome** no one can be made better off without making someone else worse off

## Definition (Pareto Optimality)

- Strategy profile  $s$  **Pareto dominates** strategy profile  $s'$  if:
  - $u_i(s) \geq u_i(s')$  for every player  $i \in N$ , and
  - there exists an agent  $j$  such that  $u_j(s) > u_j(s')$
- Strategy profile  $s$  is **Pareto optimal** if there is no profile  $s'$  that Pareto dominates  $s$

# The Coordination Game

- A country with no traffic rules
- Two cars on the road, driving towards each other
- They have to decide what side of the road to take
- If they choose the same side, all is well
- If they choose different sides, they crash into each other



# The Coordination Game

What are the Pareto optimal outcomes?

	Left	Right
Left	(1,1)	(0,0)
Right	(0,0)	(1,1)

Pareto dominated by (Left, Left) and (Right, Right)

Pareto optimal



# More on Pareto Optimality

- Pareto optimal outcomes always exist
- They are not always unique (there can be more than one)
- Pareto optimality does not necessarily imply that outcomes are fair

